

Your grade: 100%**Next item →**Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.**1 / 1 point**

1. True/False: Suppose you learn a word embedding for a vocabulary of 60000 words. Then the embedding vectors could be 60000 dimensional, so as to capture the full range of variation and meaning in those words.

 False True**Expand****Correct**

No, the dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. What is t-SNE?

1 / 1 point

- A supervised learning algorithm for learning word embeddings
- A linear transformation that allows us to solve analogies on word vectors
- A non-linear dimensionality reduction technique
- An open-source sequence modeling library

Expand**Correct**

Yes

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

1 / 1 point

x (input text)	y (happy?)
Having a great time!	1
I'm sad it's raining.	0
I'm feeling awesome!	1

Even if the word "wonderful" does not appear in your small training set, what label might be reasonably expected for the input text "I feel wonderful!"?

 y=1 y=0**Expand****Correct**

Yes, word vectors empower your model with an incredible ability to generalize. The vector for "wonderful" would contain a negative/unhappy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply)

1 / 1 point $e_{man} - e_{woman} \approx e_{king} - e_{queen}$ **Correct**

The order of words is correct in this analogy.

 $e_{man} - e_{woman} \approx e_{queen} - e_{king}$ $e_{man} - e_{king} \approx e_{queen} - e_{woman}$ $e_{man} - e_{king} \approx e_{woman} - e_{queen}$ **Typesetting math: 100%**

The order of words is correct in this analogy.

Expand**Correct**

Great, you got all the right answers.

5. True/False: The most computationally efficient formula for Python to get the embedding of word 1021, if C is an embedding matrix, and o_{1021} is a one-hot vector corresponding to word 1021, is $C^T * o_{1021}$.

1 / 1 point True False**Expand****Correct**

It is computationally wasteful because the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we create an artificial task of estimating $P(\text{target} | \text{context})$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.

1 / 1 point False True**Expand****Correct**

and are chosen from the training set to be nearby words.

8. Suppose you have a 10000 word vocabulary, and are learning 100-dimensional word embeddings. The word2vec model uses the following softmax function:

1 / 1 point

$$P(t | c) = \frac{e^{t^T e_c}}{\sum_{i=1}^{10000} e^{t^T e_i}}$$

True/False: After training, we should expect θ_t to be very close to e_c when t and c are the same word.

 True False**Expand****Correct**

To review this concept watch the lecture.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

1 / 1 point

$$\min \sum_{i=1}^{10000} \sum_{j=1}^{10000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - \log X_{ij})^2$$

Which of these statements are correct? Check all that apply.

 X_{ij} is the number of times word j appears in the context of word i .**Correct** Theoretically, the weighting function $f(\cdot)$ must satisfy $f(0) = 0$ **Correct** θ_i and e_j should be initialized randomly at the beginning of training.**Correct**

Typesetting math: 100% θ_i should be initialized to 0 at the beginning of training.

Expand**Correct**

Great, you got all the right answers.

10. You have trained word embeddings using a text dataset of s_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of s_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

1 / 1 point $s_1 \ll s_2$ $s_1 \gg s_2$ **Expand****Correct**

s_1 should transfer to s_2